

When the OS gets in the way

(and what you can do about it)

Mark Price
@epickrram

LMAX Exchange

Linux

When the **OS** gets in the way

(and what you can do about it)

Mark Price
@epickrram

LMAX Exchange

It's not the OS's fault

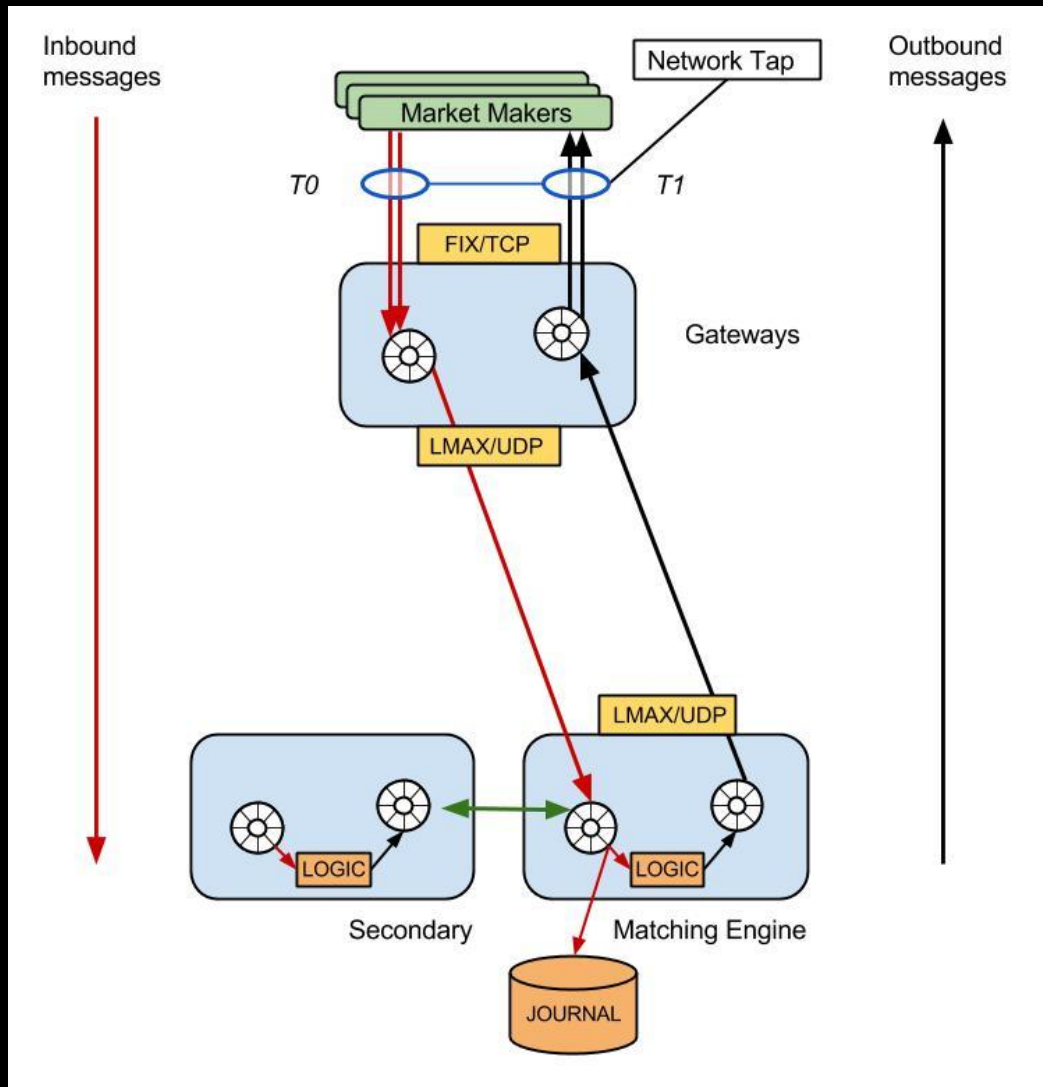
- Linux is an excellent general-purpose OS
- Many target platforms
- Scheduling is actually fairly complicated
- Low-latency is a special use-case
- We need to provide some hints

Why should I care?

Useful in some scenarios

- Low latency applications
- Response times $< 1\text{ms}$
- Compute-intensive workloads
- Long-running jobs

A real-world scenario: LMAX



System

$$\text{Latency} = T1 - T0$$

Before tuning:

250us / 10+ms

After tuning:

80us / <1ms

(mean / max)

Jitter

- *“slight irregular movement, variation, or unsteadiness, especially in an electrical signal or electronic device”*
- Variation in response time latency
- Long-tail in response time

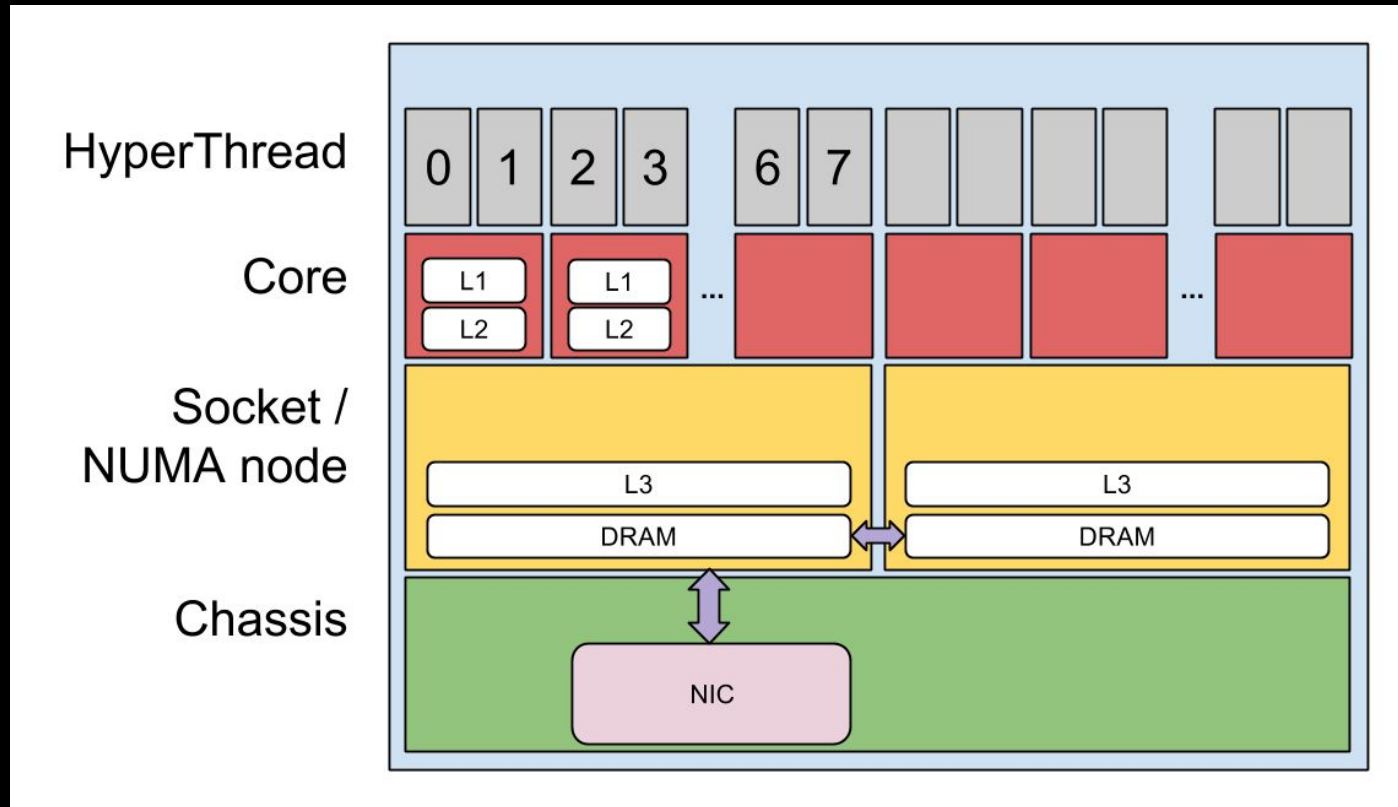
Dealing with it

- First take care of the low-hanging fruit
 - e.g. Garbage collection (gc-free / Zing)
 - e.g. Slow I/O
- Once response times are $< 10\text{ms}$ the fun begins
- Make sure your code is running!

Measure first

- Need to validate changes are good
- End-to-end tests
- Using realistic load
- Change one thing and observe
- A refresher...

Modern hardware layout




Multi-tasking

- $\text{num(tasks)} > \text{num(HyperThreads)}$
- OS must share out hardware resources
- Clever? Dumb? Fast? Slow?
- Fair...


Linux CFS

- Completely Fair Scheduler
- Maintains a task 'queue' per HT
- Runs the task with the lowest runtime
- Updates task runtime after execution
- Higher priority implies longer execution time
- Tasks are load-balanced across HTs

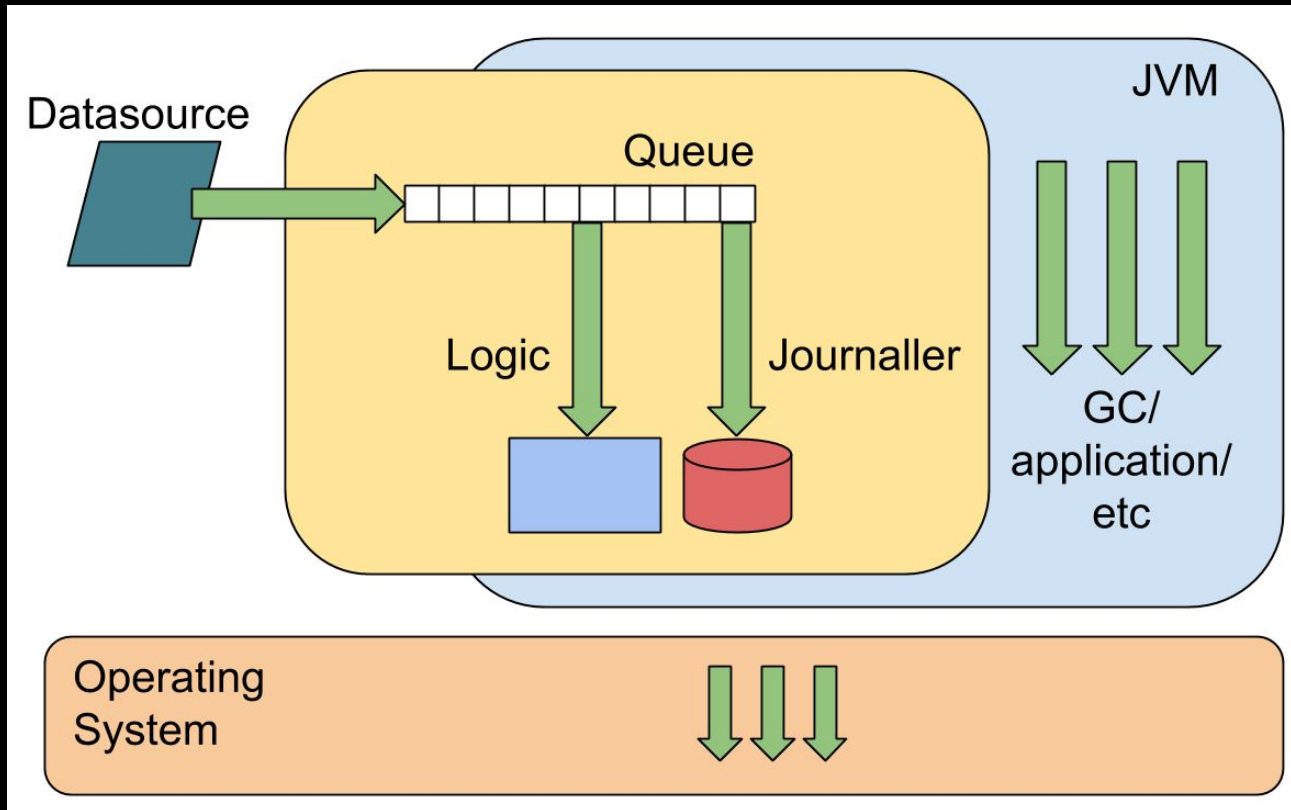
An example application ...




... running on a language runtime



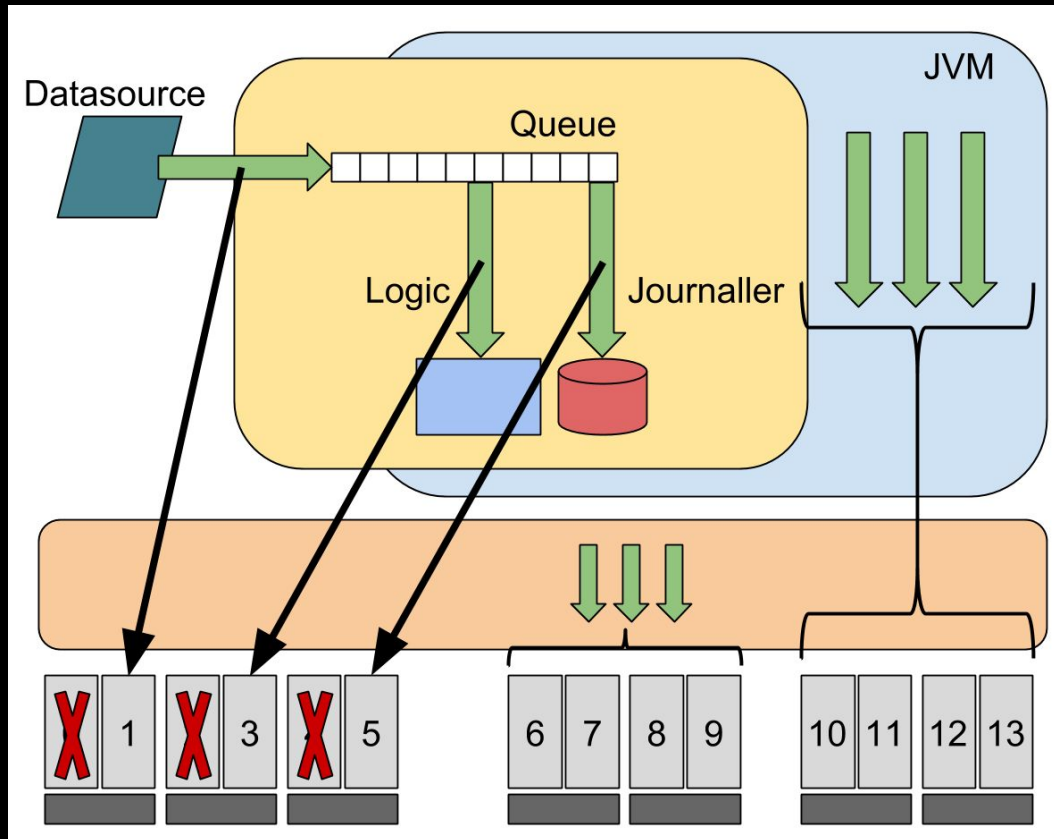
... running on an operating system



Optimise for locality - PCI/memory



Target deployment



How do I start?

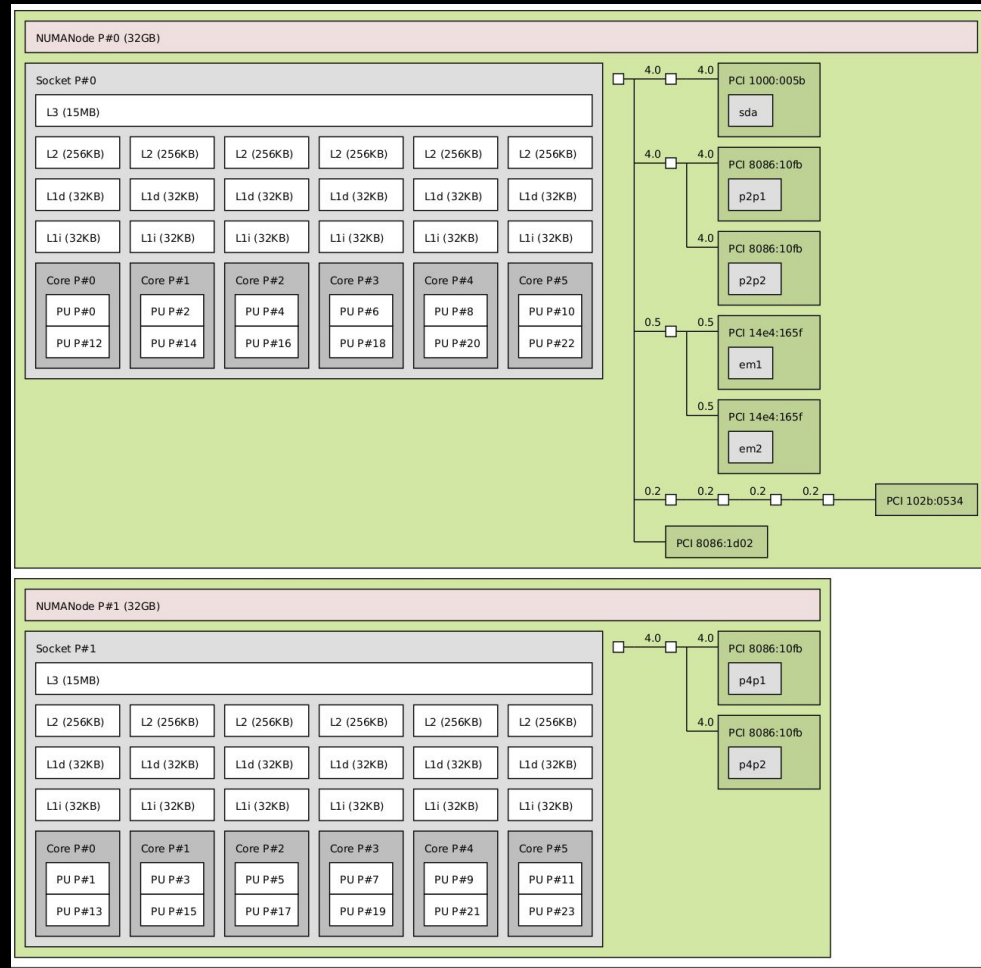
Start with the metal

- BIOS settings for maximum performance
- That's a whole other talk...

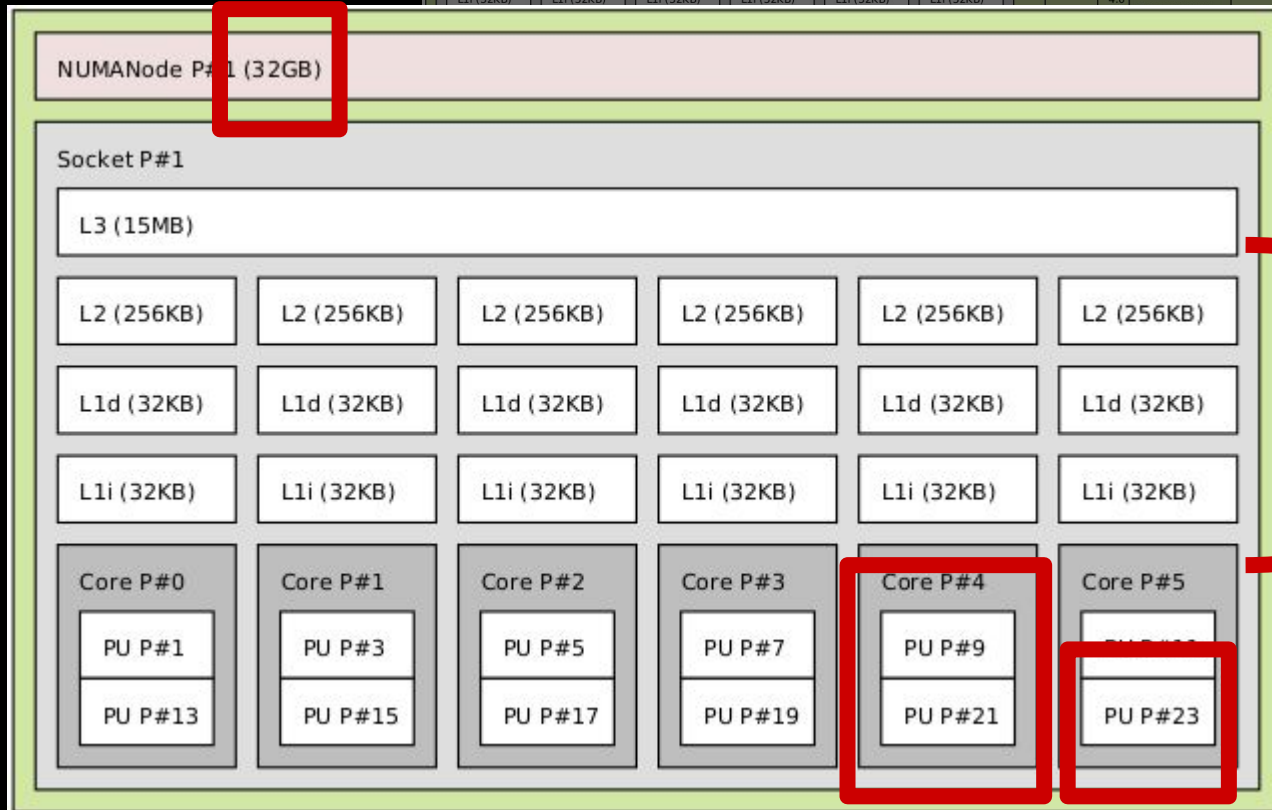

Discover what's available

- *lstopo* is a useful tool for looking at hardware
- Provided by the hwloc package
- Displays:
 - HyperThreads
 - Physical cores
 - NUMA nodes
 - PCI locality

Istopo



Istopo



NUMA-local
RAM

Caches

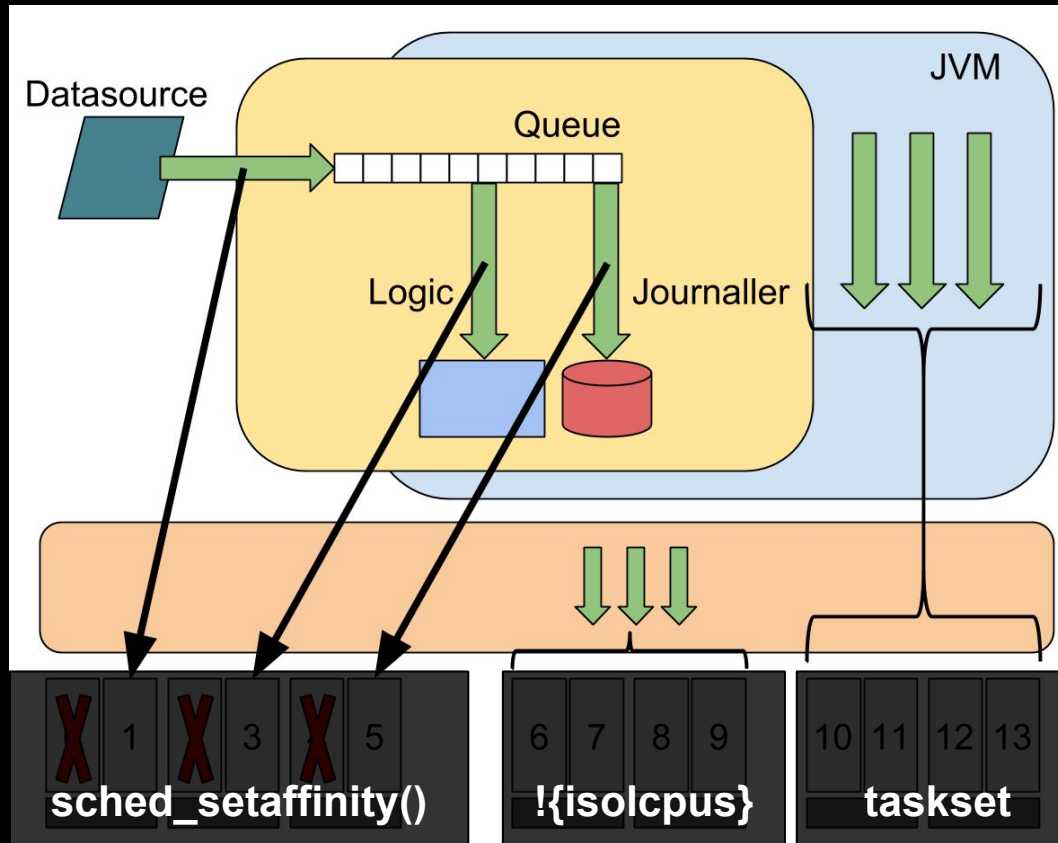
Core

HyperThread


Reserve & use specific resource

- Use *isolcpus* to reserve cpu resource
- kernel boot parameter
- `isolcpus=0-5,10-13`
- Use *taskset* to pin your application to cpus:
- `taskset -c 10-13 java ...`
- Set affinity of hot threads:
- *`sched_setaffinity(...)`*

Deploy the application



You have no load-balancer



A solution: cpusets

- Create hierarchical sets of reserved resource
- CPU, memory
- Userland tools: cset (SUSE)

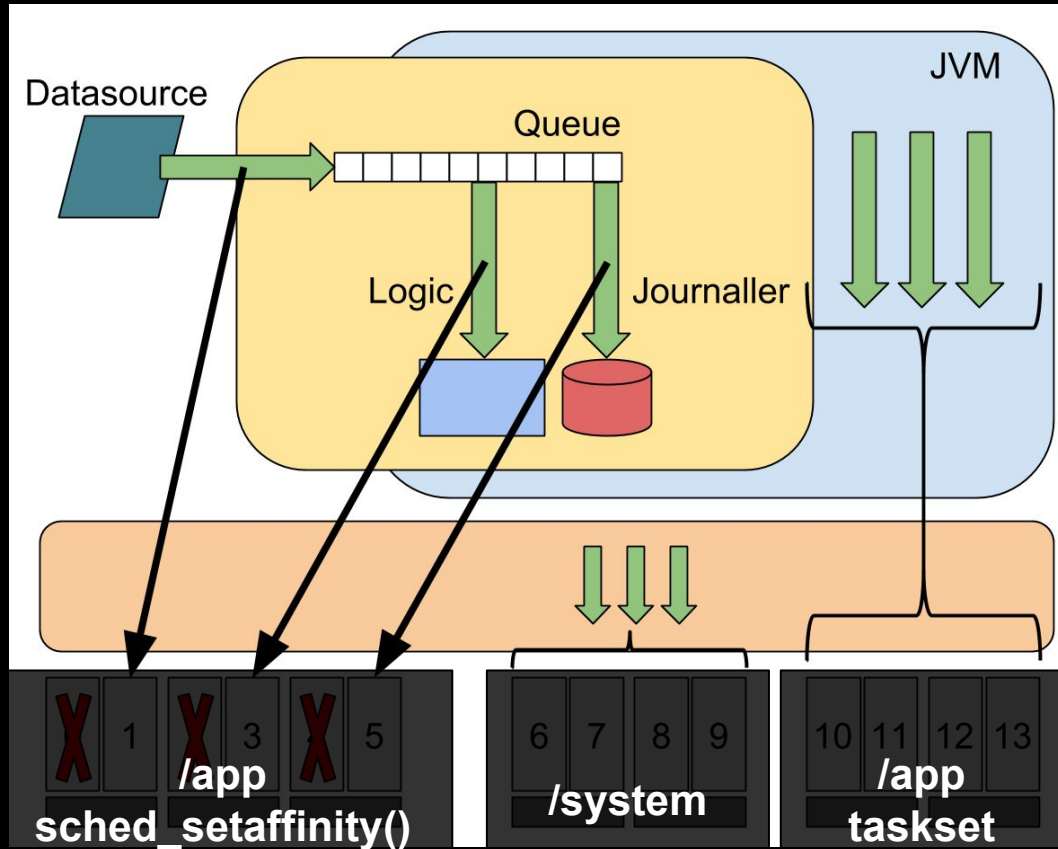
Isolate OS processes

- `cset set --set=/system --cpu=6-9`
 - *create a cpuset with cpus 6-9*
 - *create it at the path /system*
- `cset proc --move --from-set=/ --to-set=/system`
 - *move all processes from / to /system*
 - *-k => move unbound kernel threads*
 - *--threads => move child threads*
 - *--force => erm... force*

Run the application

- `cset set --cpu=0-5,10-13 --set=/app`
- `cset proc --exec /app taskset -cp 10-13 java ...`
 - *start a process in the /app cpuset*
 - *run the program on cpus 10-13*
- `sched_setaffinity()` to pin the hot threads to cpus 1,3,5

Isolated threads



No more jitter?

perf_events

- Sampling tracer
- Static/dynamic trace points
- Very low overhead
- A good starting point for digging deeper
- *perf list* to view available trace points
- network, file-system, scheduler, etc

What's happening CPU?

- `perf record -e "sched:sched_switch" -C 3`
 - Sample task switches on CPU 3
- `perf report` (best for multiple events)
- `perf script` (best for single events)

Rogue process

java

36049 [003] 3011858.780856: sched:sched_switch: java:
36049 [110] R ==> kworker/3:1:13991 [120]

kworker/3:1

13991 [003] 3011858.780861: sched:sched_switch:
kworker/3:1:13991 [120] S ==> java:36049 [110]

ftrace

- Function tracer
- Static/dynamic trace points
- Higher overhead
- But captures everything
- Can provide function graphs
- trace-cmd is the usable front-end

So what is that kernel thread doing?

- `trace-cmd record -P <pid> -p function_graph`
 - Trace functions called by process <pid>
- `trace-cmd report`
 - Display captured trace data

Some things can't be deferred

```
kworker/3:1-13991 [003] 3013287.180771: funcgraph_entry:      | process_one_work() {  
kworker/3:1-13991 [003] 3013287.180772: funcgraph_entry:      |     cache_reap() {  
kworker/3:1-13991 [003] 3013287.180772: funcgraph_entry: 0.137 us |         mutex_trylock();  
kworker/3:1-13991 [003] 3013287.180772: funcgraph_entry: 0.289 us |         drain_array();  
kworker/3:1-13991 [003] 3013287.180773: funcgraph_entry: 0.040 us |         _cond_resched();  
.....  
.....  
kworker/3:1-13991 [003] 3013287.180859: funcgraph_exit: +86.735 us | }
```



+86.735 us

Things to look out for

- `cache_reap()` - SLAB allocator
- `vmstat_update()` - kernel stats
- other workqueue events
 - `perf record -e "workqueue:*" -C 3`
- Interrupts - set affinity in `/proc/irq`
- Timer ticks - tickless mode
- CPU governor - set to performance
 - `/sys/devices/system/cpu/cpuN/cpufreq/scaling_governor`


Some numbers

- Inter-thread latency is a good proxy
- 2 busy-spinning threads passing a message
- Time taken between producer & consumer
- Record times over several seconds
- Compare tuned/untuned


Results

== Latency (ns) ==	untuned	tuned
mean	466	216
min	200	128
50.00%	464	208
90.00%	608	288
99.00%	768	336
99.90%	992	544
99.99%	2432	1664
max	69632	69632

tuned vs untuned




tuned vs untuned (log scale)



Results (loaded system)

== Latency (ns) ==	untuned	tuned
mean	545	332
min	144	216
50.00%	464	336
90.00%	544	352
99.00%	736	448
99.90%	2944	544
99.99%	294913	704
max	884739	36864

tuned vs untuned (loaded system)



Summary

- Select threads that need access to CPU
- Isolate CPUs from the OS
- Pin important threads to isolated CPUs
- Don't forget interrupts
- There will be more things...
- Always test assumptions!
- Run validation tests to ensure tunings are as expected

Thank you

- lmax.com/blog/staff-blogs/
- epickrram.blogspot.com
- github.com/epickrram/perf-workshop
- [@epickrram](#)